

A clear, public submission.

Repository and judge guide - checked 5 October 2026. The product repository is public and the Linux-hosted app is live. Verified results and remaining limits are recorded below.

What the official guide requires

- **Public GitHub repository.** Private repositories are not accepted.
- **Complete permitted code and runtime files.** Include dependencies, rights notices and instructions explaining the idea, setup and operation.
- **Source, tool and license record.** Disclose AI services, models, open-source components and data/asset origins.
- **No secrets or beneficiary data.** Personal research, accounts and real recordings stay outside the public release.
- **Working live product link.** Test the actual hosted functions before final submission.

The supplied final judging rubric

Criterion	Weight
Technical solution and AI use	25%
Reliability and scholarly safety	15%
Innovation and added value	15%
Beneficiary experience, communication and accessibility	10%
Benefit against the track success criterion	20%
Operational realism and continuation	10%
Presentation clarity and verifiability	5%

Weights come from the supplied final dashboard. These are criteria, not a predicted DarsLoop score. Feature counts and local tests do not prove student-learning benefit.

Other submission materials

The observed final form asks for project details, the accepted path, one presentation file (PDF/PPT/PPTX, up to 10 MB), a video link of no more than two minutes, GitHub URL and live-demo URL. Keep the planned ten-slide deck within the registration limit. Submit through the official portal and keep the receipt.

Published guide deadline: 6 October 2026, 23:59 Riyadh / 7 October, 01:59 Pakistan. Check actual later notices; do not assume an extension.

Basis: archived official Participant Guide, PDF positions 29-32; supplied portal/rubric; current [official terms](#), sections 8-10. The public guide fetch failed during this recheck. Muse is a secondary checklist; official requirements take priority.

A short path for judges.

A prepared fictional example works without account setup or AI keys. Fresh audio processing is a separate, connected-service workflow.

1. Open the prepared example

[Live fictional example](#) | [Public GitHub repository](#)

Install Node.js 24 or later with npm. Download/clone the public repository and open a terminal in the folder containing package.json. The actual GitHub clone was checked with Node 26.4.0 on macOS.

```
npm ci

npm run build

npm start
```

Open **<http://127.0.0.1:3000/example>**. No environment file is needed. The prepared example does not call FFmpeg or an AI provider. It is labelled authored fictional material; it is not fresh generation.

2. Follow the lesson evidence

- Open a note and play its timestamp. Compare the captured quote with the recording.
- In **Ask this lesson**, ask about missing a lesson and follow its supported passage. Ask an absent topic and check the scope response.
- Ask **Is vaping halal?** and check teacher referral. The app must not issue a personal ruling.
- Try a quiz and a flashcard, then inspect the linked audio. Example practice resets when leaving.
- Open Catch me up, Teacher's terms and Transcript. Keep captured wording separate from AI explanation.

3. Test fresh audio when services are configured

Install FFmpeg and FFprobe. Copy the empty .env.example to .env.local only if it does not already exist. Add your own Groq/Google keys privately; restart the app and run **npm run worker** in a second terminal. Open /learn and upload the bundled fictional demo.mp3 with the permission/fictional confirmations.

Choose Quick, Balanced or Detailed notes. Wait for real worker completion, then check transcript times, uncertain passages, notes, quizzes and cards. Without keys, a fresh upload cannot finish AI processing. Prepared output is never substituted for it.

4. Reproduce the checks

npm run typecheck and **npm test** check code and units; the full tests need FFmpeg/FFprobe. **npm run test:api** starts an isolated temporary server and removes its data. **npm run test:example** checks the running example. Cloud/provider checks are optional and require private keys plus service quota.

Public hosting uses Supabase authentication/storage and a web+worker Docker service. See README.md, DEPLOYMENT.md and docs/JUDGE_GUIDE.md for the full setup. Local SQLite mode is loopback-only.

Publish the product only.

The release uses an explicit file allowlist and a private-value scan. A .gitignore file alone is not considered proof that a package is safe.

Files judges receive

File or folder	Purpose
README.md	Quick start, prerequisites, AI setup and known limits
docs/JUDGE_GUIDE.md / ARCHITECTURE.md	Walkthrough, data flow and implemented boundaries
DEPLOYMENT.md / docs/QA.md	Hosting/recovery instructions and measured verification scope
SOURCES.md / THIRD_PARTY_NOTICES.md / LICENSE	Source, tool, model, asset and license disclosure
.env.example / package-lock.json	Empty connection template and pinned dependency versions
src / scripts / tests / supabase	Application, curated runtime/check scripts, tests and safe migrations
public / fixtures	Original illustrations and labelled fictional lesson material

Files judges do not receive

Personal documents, Muse drafts, private research inputs, dashboard screenshots/PDFs, conversation history, passwords, API-key values, account credentials, databases, real student/class recordings, raw private test reports, node_modules and build caches.

What has actually been checked

The public GitHub clone installed and built without keys, with 87 unit tests and 18 isolated API checks. Railway passed 14 HTTPS example checks, 26 fresh authenticated study assertions and 23 direct-upload assertions. A 38 MB fictional source resumed across deployment and produced actual notes/practice. Container peak was 292 MB under a 1 GB limit, with no OOM kills. Full scopes are in docs/QA.md.

What still needs finishing

Signup email/social consent, in-flight worker crash recovery, genuine starting-version/reused-work disclosure, final deck/video and submission receipt. No backdated commit history, learner pilot or broad accuracy/learning claim is invented. Trial credits and provider quotas remain finite; desktop monitoring requires the desktop to be available. See DEPLOYMENT.md for the dated hosting checklist.

Browser lock can interrupt recording. Teacher handoff is optional and is not authenticated approval. Lesson tests are not official curriculum exams. Quotas can interrupt fresh processing. These practical limits remain visible to judges.

Official reference links

[Official terms](#) | [Participant Guide](#) | [Scientific Framework](#)